Advanced Edge-Native Architectures for Autonomous Rapid Crisis Detection and Management

The proliferation of smart city infrastructure has precipitated a critical evolution in how municipal authorities, emergency responders, and private enterprises approach disaster management. Historically, automated crisis detection has been constrained by centralized, cloud-tethered architectures. In these legacy models, distributed sensors—ranging from closed-circuit television (CCTV) to environmental monitors—stream raw data to distant cloud servers for processing, analysis, and alert generation. While this paradigm leverages immense computational resources, it introduces fatal vulnerabilities during high-stress scenarios. Natural disasters, such as severe earthquakes or catastrophic wildfires, and localized emergencies, such as structural fires or mass stampedes, frequently compromise primary telecommunications infrastructure. When the network backhaul is severed, cloud-dependent emergency systems fail precisely at the moment they are most desperately required, rendering sensor networks blind and automated response protocols inert.[1]

To overcome the fragility of centralized topologies, the industry is undergoing a paradigm shift toward edge-native, offline-first architectures. By relocating the computational workload directly to the geographical boundary of the data generation—deploying lightweight artificial intelligence (AI) inference models onto highly constrained but capable microcomputers—systems can achieve total operational autonomy. This comprehensive report investigates the design, algorithmic optimization, and deployment strategy of a multi-modal, edge-driven rapid crisis detection framework. This infrastructure focuses primarily on the immediate, localized detection of earthquakes, fires, and crowd stampedes, while retaining a highly modular software architecture to accommodate future hazard integrations, such as individual medical emergencies and structural entrapment.[3]

Furthermore, a holistic response framework cannot exist solely as a passive detection grid. It requires robust integration with autonomous AI coordination agents to dispatch resources, a geofenced mobile application to interface directly with affected civilians, and a centralized web dashboard compliant with Next Generation 9-1-1 (NG9-1-1) protocols for authoritative oversight. Through the synthesis of edge computing, specialized deep learning architectures, and human-centered crisis communication design, this report delineates a state-of-the-art approach to modern disaster resilience.

Hardware Infrastructure and Edge Compute Processing Strategies

The foundational layer of the crisis detection system is governed by the physical hardware

deployed in the field. The selection of single-board computers (SBCs) dictates the maximum allowable complexity of the neural networks, the latency of real-time inference, and the fundamental power autonomy of the node. The architectural topology relies primarily on two competing but complementary hardware ecosystems: the general-purpose ARM architecture of the Raspberry Pi and the GPU-accelerated architecture of the NVIDIA Jetson series.[6]

Compute Node Selection: Raspberry Pi 5 vs. NVIDIA Jetson

Deploying complex computer vision, temporal sequence modeling, and large-scale sensor aggregation at the edge requires a delicate balance between computational throughput, thermal management, and financial scalability. Computational throughput in AI workloads is standardly measured in Tera Operations Per Second (TOPS).

The NVIDIA Jetson architecture represents the pinnacle of AI-accelerated edge processing due to its native integration of highly parallelized GPU architectures (such as Maxwell, Ampere, or Ada Lovelace) featuring hundreds of dedicated Compute Unified Device Architecture (CUDA) cores.[6] The legacy Jetson Nano, even with its aging 128-core Maxwell GPU, provides an approximate 22x advantage in AI performance over standard CPU-bound boards.[7] This CUDA optimization allows deep learning frameworks, including TensorFlow, PyTorch, and OpenCV, to execute hardware-accelerated tensor operations. In direct benchmark comparisons for real-time computer vision tasks, such as running the ResNet-50 architecture, the Jetson Nano processes up to 36 frames per second (FPS), whereas unaccelerated legacy Raspberry Pi models struggle to achieve 1.4 FPS.[7] For highly intensive, modern workloads involving multi-stream video analytics, complex crowd optical flow, and on-board generative AI triage agents, the newer Jetson Orin Nano Super provides a massive 67 TOPS, representing the optimal hardware deployment for high-density environments.[6]

Conversely, the Raspberry Pi 5 dominates the domain of general-purpose computing, lightweight server orchestration, and dense sensor Input/Output (I/O) aggregation.[6] While the Pi 5 lacks a dedicated CUDA-enabled GPU, its quad-core ARM Cortex-A76 architecture delivers exceptional single-core and multi-core performance, scoring approximately 800 and 2,500 respectively on Geekbench 5 tests—vastly outperforming the CPU benchmarks of the legacy Jetson Nano.[9] For systems requiring real-time computer vision without the budget or thermal envelope of an NVIDIA device, the Raspberry Pi 5 can achieve real-time performance on lightweight, highly quantized models (such as YOLOv5n or MediaPipe) when coupled with specialized Neural Processing Unit (NPU) accelerators like the Hailo-8 or the Coral Tensor Processing Unit (TPU).[8]

Hardware Platform	Primary Architecture Focus	AI Throughput Capabilities	Optimal Crisis Detection Workloads

Raspberry Pi 5	Quad-core ARM Cortex-A76	CPU-bound (unless supplemented by external TPU)	Environmental sensor aggregation, lightweight server hosting, TinyML seismic evaluation, and MediaPipe skeletal tracking.[6]
Jetson Nano (Legacy)	128-core Maxwell GPU	0.5 TOPS	Single-stream YOLO visual inference, basic fire/smoke detection, tight budget constraints.[6]
Jetson Orin Nano	Ampere architecture GPU	Up to 67 TOPS	Multi-stream CCTV analysis, dense crowd/stampede prediction mapping, complex autonomous agent orchestration.[6]

Power Autonomy and Disaster Resilience Topologies

A fundamental mandate of offline-first, disaster-resilient infrastructure is its survivability during complete power grid failures.[12] An edge node relying on facility power is inherently flawed, as catastrophic events instantly disable municipal electrical grids. Thus, uninterruptible power supply (UPS) modules and renewable energy integration are critical dependencies for long-term autonomous nodes.[15]

For indoor and localized deployments of the Jetson Nano and Raspberry Pi, advanced UPS Hardware Attached on Top (HAT) modules provide intelligent, automated power management.[12] These HATs physically stack onto the GPIO headers of the SBCs and utilize arrays of 18650 lithium-ion polymer batteries. Modern UPS HATs support advanced power path management, fast battery charging capabilities (up to 3A), integrated overvoltage protection circuits, and I2C bus communication.[12] This I2C integration allows the edge server's operating system to query the Maxim fuel gauge chip in real-time, reading precise battery voltage and percentage data to initiate a graceful software shutdown if total depletion is imminent, thereby preventing fatal SD card or file system corruption during hard power loss.[12] In a disaster scenario, an appropriately configured multi-cell 18650 matrix can provide 10 to 30 hours of continuous AI inference and monitoring operations.[12]

For long-term, autonomous nodes deployed in remote or highly exposed outdoor environments—such as wildfire perimeters or expansive smart-city intersections—integration with solar photovoltaic (PV) power systems ensures theoretical perpetual operation.[18] Because edge nodes deployed in these environments are susceptible to environmental degradation, sophisticated predictive maintenance models must be deployed. By utilizing the Jetson Nano to run specialized YOLOv8 computer vision models on the node's own solar panels, the system can autonomously detect dust accumulation and solar soiling.[20] Identifying these efficiency-degrading factors allows the node to request localized maintenance before power generation drops below the threshold required to sustain the critical disaster-detection workloads.[20]

Algorithmic Frameworks for Multi-Hazard Detection

The analytical pipeline of the crisis detection system relies on highly specialized, heavily quantized AI models tailored to detect specific emergency scenarios. By processing the raw sensor and video data locally, the system entirely eliminates cloud transmission latency—a factor that drastically increases casualty rates in life-safety applications.[2]

Seismic Event Detection via Tiny Machine Learning

Earthquake Early Warning (EEW) systems typically rely on massive, macro-scale seismic networks utilizing deeply buried seismometers. However, the integration of always-on, stationary microelectromechanical systems (MEMS) accelerometers—such as the ubiquitous ADXL335—directly into localized IoT edge devices enables high-fidelity, single-station seismic monitoring.[5]

The algorithms deployed on these microcontroller and SBC nodes utilize Tiny Machine Learning (TinyML) architectures.[5] Specifically, lightweight Convolutional Neural Networks (CNNs) boasting highly compressed architectures of approximately 38,000 trainable parameters are designed to continuously process real-time tri-axial (X, Y, Z) acceleration data.[5] Seismology dictates that the non-destructive primary waves (P-waves) of an earthquake arrive seconds or tens of seconds before the destructive secondary shear waves (S-waves). By ingesting and analyzing a highly optimized 2-second waveform window, these deep learning models can successfully identify P-wave signatures before structural shaking begins.[5]

When deployed on hardware equivalent to a Raspberry Pi 5 or an Arduino Cortex M4, the inference latency of these models drops to a sub-7-millisecond threshold, achieving an overall accuracy of 97.12% and a 98% identification rate for critical P-wave segments.[5] The profound second-order implication of this extreme low-latency, edge-based processing is the mitigation of the standard "blind zone" effect. In near-field seismic events, the time required to transmit an accelerometer reading to a centralized cloud server, process the waveform, and transmit an alert back to the physical site often exceeds the arrival time of the destructive shockwave. Processing the waveform entirely locally ensures that immediate physical actuations—such as the automated severing of gas lines, the unlocking of electromagnetic emergency exits, and

the triggering of high-decibel audible alarms—occur instantaneously, significantly reducing secondary disaster risks.[5]

Fire, Smoke, and Chemical Hazard Optical Verification

Traditional fire detection infrastructure relies heavily on primitive particulate scattering (standard smoke detectors) or rudimentary thermal thresholds. These analog systems are highly susceptible to false positives generated by ambient dust, industrial steam vapors, or non-threatening temperature fluctuations, leading to debilitating alert fatigue.[24] The proposed advanced architecture utilizes a multimodal AI approach, fusing visual computer vision models with parallel thermal and chemical sensor telemetry to achieve near-perfect reliability.[25]

For the primary visual detection vector, the system deploys lightweight versions of the "You Only Look Once" (YOLO) object detection architecture, specifically the "nano" (n) and "tiny" (t) variants like YOLOv5n, YOLOv8n, and YOLO11n, which are explicitly engineered for CPU and low-power GPU edge inference.[4] Recent empirical research evaluating these specific models for wildfire and structural fire detection on edge hardware highlights a highly counter-intuitive finding regarding the Energy-Delay Product (EDP).[4] While YOLO11n represents a newer, more sophisticated architecture, the older YOLOv5n—featuring only 1.8 million parameters and 5.1 billion Floating Point Operations Per Second (FLOPs)—performs significantly faster on CPU-bound edge devices lacking hardware acceleration.[4] In testing, YOLOv5n processed high-resolution images at approximately 6.1 FPS, nearly double the 3.3 FPS achieved by YOLO11n.[4] Because the inference time is drastically reduced, the total average power consumed per detection (approximately 6886 mW to 6924 mW) yields the best normalized EDP, making YOLOv5n the optimal choice for battery-constrained nodes operating on solar or UPS power.[4] Through multi-stage Transfer Learning (TL) on localized source datasets like the Flame and Smoke Detection Dataset (FASDD), the edge model achieves an impressive mean Average Precision (mAP@0.5) of 79.2%, with highly accurate bounding box confidence scores for both dynamic flames and faint smoke plumes.[4]

However, relying solely on visual AI can still yield false positives under challenging lighting conditions or complex architectural geometries.[25] To mitigate this, the system cross-references visual YOLO detections with concurrent data streams from thermal imaging cameras and ambient chemical sensors (detecting COx and NOx spikes).[24] The edge hub executes a secondary CNN that analyzes the geometric density of the visual smoke plume while simultaneously evaluating thermal heat-signature gradients and localized oxygen reduction.[25] A critical alert is only escalated to the autonomous coordination agents if a true multimodal signature is mathematically confirmed, effectively eliminating false alarms and preserving critical emergency response resources.[25]

Stampede Prediction and Deep Crowd Density Analytics

Mass gatherings in urban centers, transport hubs, and religious events present severe logistical and safety challenges. Rapid fluctuations in crowd density can quickly escalate into fatal

stampedes, asphyxiation events, and crowd crushes, often developing into highly lethal scenarios within minutes.[31] Conventional bounding-box object detection algorithms (like standard YOLO implementations) struggle severely in ultra-dense crowd environments due to massive human-to-human occlusion, where identifying individual limbs and heads becomes mathematically intractable.[34]

To overcome this scale and occlusion challenge, the architecture integrates a highly specialized hybrid CNN-LSTM (Convolutional Neural Network combined with a Long Short-Term Memory network) coupled with dense Farneback optical flow computation.[3] Rather than attempting to draw bounding boxes around discrete individuals, the Farneback optical flow algorithm computes motion-centric indicators across the entire temporal video frame.[3] It measures pixel-level vector fields to determine global crowd dynamics, extracting highly granular data on flow acceleration, vector divergence, and localized flow entropy.[3] The CNN layer acts as a lightweight spatial feature encoder, while the LSTM layer captures the short-term temporal coherence of the crowd's kinetic movement over consecutive frames.[3]

This paradigm fundamentally shifts crowd management from a reactive posture to a highly predictive one. The system categorizes crowd risk into a precise four-class taxonomy: normal, moderate, dense, and risky.[3] By analyzing 30-second sliding windows of optical flow and crowd density data (measured in persons per square meter), the LSTM model can accurately forecast high-risk zones up to two minutes before a critical, life-threatening surge materializes.[33] When the spatial grid density approaches a critical threshold of 4 persons/m^2, or when sudden vector divergence indicates explosive panic dispersal, the edge node autonomously triggers early intervention protocols.[3] Validated on rigorous domain-shift benchmarks such as the UMN anomaly dataset, this hybrid architecture has demonstrated an experimental accuracy of 99.75% in distinguishing between safe high-density movement and hazardous stampede conditions.[3]

Individual Entrapment and Medical Fainting Detection

While macroscopic hazard detection (fires, earthquakes, stampedes) protects large populations, a comprehensive rapid crisis system must also scale down to detect individual, localized emergencies. In the future expansion of this system, detecting a single individual fainting in a corridor or becoming trapped beneath structural debris following a collapse is paramount.

For the real-time anomaly detection of human posture indicating a medical emergency (e.g., sudden collapse, seizure, or fainting), the edge nodes leverage the highly optimized MediaPipe Pose framework developed by Google.[11] MediaPipe provides lightweight, real-time human pose estimation by tracking 33 three-dimensional skeletal landmarks on the human body.[11] By running kinematic analysis on the geometry, velocity, and relative positioning of these keypoints—such as observing the sudden, unbraced acceleration of the cranial nodes toward the floor plane without a corresponding forward walking vector—the system accurately and anonymously classifies a fainting event.[11] Because MediaPipe is ruthlessly optimized for

multimedia processing, a standard Raspberry Pi 5 can run the pose estimation pipeline smoothly in real-time, while the GPU-accelerated Jetson Nano can achieve frame rates up to three times higher than legacy Pi hardware, ensuring instantaneous medical alerts.[11]

In highly destructive scenarios involving complete structural collapse, standard optical cameras are rendered entirely blind by dust, rubble, and darkness. The framework addresses human entrapment in these environments through multi-modal alternative sensing:

1. **Acoustic AI Detection**: Unmanned Aerial Vehicles (UAVs) acting as mobile edge nodes, or surviving stationary nodes, utilize directional microphones to capture ambient audio. Deep learning CNNs are trained specifically on acoustic signatures to filter out mechanical drone motor noise and environmental interference, mathematically isolating distinct human distress signals such as screams, shouts, or rhythmic tapping against pipes.[40]
2. **Ultra-Wideband (UWB) and Penetrating Radar**: Specialized edge nodes deployed into the rubble can utilize UWB radar technologies to detect microscopic movements beneath concrete debris. These radar algorithms can isolate the rhythmic patterns of human respiration and heartbeat frequencies, bypassing visual occlusion entirely and providing precise localization of trapped survivors to rescue personnel.[43]

Autonomous AI Agents and Emergency Coordination Workflows

To process the overwhelming deluge of sensor telemetry, optical alerts, and user-submitted mobile reports during a mass casualty event, the system employs a layer of autonomous AI agents. The sheer volume of simultaneous incoming alerts during a macroscopic disaster can easily overwhelm human operators, leading to dispatch paralysis.[46]

Multi-Agent Swarm Architecture

The logical, decision-making brain of the response system is constructed using advanced Large Language Model (LLM) orchestration frameworks, specifically LangChain and LangGraph.[49] While earlier frameworks like AutoGPT operate on fully autonomous, recursive, self-prompting loops that frequently suffer from fatal errors, context degradation, or infinite logic loops during complex tasks, LangChain and LangGraph provide a much safer alternative.[49] These libraries allow developers to create deterministic, stateful, cyclical graphs with strict programmatic guardrails, making them exponentially more robust for mission-critical business continuity and disaster response engineering.[49]

The AI backend is architected as a collaborative "swarm" of specialized agents:

1. **The Emergency Coordinator (EC)**: This agent acts as the primary triage hub, functioning analogously to a highly advanced 911 dispatcher. It continuously parses incoming JSON payloads from the edge sensors and natural language text reports originating from the

civilian mobile application.[51]

2. **Specialist Sub-Agents**: Based on the exact nature of the EC's triage, sub-tasks are delegated to highly specialized LLM agents possessing specific tool access. For example, a localized structural collapse sensor alert triggers the *Disaster Response Expert* agent, while a civilian user reporting sudden chest pains via the app routes directly to the *Medical Evacuation Specialist* agent.[51]

Automated Dispatch and Interventional Communication

The integration of these AI agents with external telecommunication APIs—such as Twilio for Programmable Messaging (SMS/RCS) or Slack/Discord webhooks for internal team routing—enables fully autonomous, real-time crisis communication.[54] When an edge node detects a high-risk stampede scenario developing in a specific hallway, the EC agent automatically synthesizes a precise warning message and utilizes the Twilio tool integration to dispatch an SMS to all registered mobile devices currently located within that specific geofenced grid.[51]

Furthermore, these agents contain embedded logic loops designed to detect immediate, severe threats to life. By utilizing programmatic severity scoring algorithms—such as the SAFE-T (Safety Alert for Emergency Triage) logic model—an agent can mathematically calculate real-time risk based on user inputs.[54] If a calculated severity threshold (e.g., >= 0.9) is breached, the agent immediately suspends all normal conversational routines, triggers a high-priority interrupt, alerts human emergency operators via instant webhooks, and dispatches critical survival protocols directly to the user's device.[54] This human-in-the-loop framework ensures that AI drastically accelerates the cognitive processing and triage of the disaster, while keeping human authorities firmly in command of macro-level strategic decisions.[46]

Edge-Native Networking and Offline-First Resilience

A highly decentralized physical deployment requires equally robust logical orchestration. Standard cloud-dependent IoT systems rely on continuous TCP/IP connections to centralized servers; if a fire severs the primary fiber backhaul, the entire system fails catastrophically.[14] To ensure absolute resilience, the system embraces a strict "offline-first" architectural paradigm, guaranteeing that devices continue to execute inference and automated control loops regardless of network status.[14]

Kubernetes Orchestration at the Edge

To manage the deployment, continuous updating, and load balancing of the various AI containers (e.g., the YOLO visual container, the TinyML seismic container, the local database) across hundreds of distributed Raspberry Pi and Jetson devices, the system utilizes lightweight Kubernetes distributions.[61] Full-scale enterprise Kubernetes is far too resource-heavy for SBCs;

therefore, specialized edge distributions like K3s, MicroK8s, and OpenYurt are deployed.[61]

K3s offers a dramatically reduced memory footprint by stripping out legacy, cloud-specific drivers while providing robust, standard-compliant container orchestration, making it a de facto standard for resource-constrained edge nodes.[61] Furthermore, platforms like OpenYurt, an extension of Kubernetes explicitly designed for edge computing, utilize local caching mechanisms to ensure that edge nodes maintain their operational autonomy and lifecycle management capabilities even when disconnected from the central control plane.[61]

Disconnected Telemetry and High-Availability Mesh Networking

In this offline-first topology, the edge nodes do not rely on an external cloud to process or act upon alerts.[14] Instead, each localized cluster acts as a self-contained intelligence hub. If an earthquake destroys the cellular towers and primary internet connections, the nodes immediately fall back to a localized mesh network, utilizing Long Range (LoRa) radio frequencies or specialized Wi-Fi ad-hoc meshes.[65]

Within this mesh architecture, a designated high-availability node (e.g., a heavily secured Jetson Orin Nano operating on dedicated solar/UPS power) serves as the local aggregator. Alerts generated by individual smart cameras and sensors are synchronized across the localized mesh using an on-premise MQTT broker secured over TLS.[25] This closed-loop configuration ensures that local physical responses—such as activating audible sirens, unlocking access-control doors, and reversing HVAC ventilation systems for smoke extraction—are executed autonomously by the edge network without a single data packet ever needing to route to an external cloud provider.[14]

The Citizen-Facing Mobile Application: Localization and Crisis UX

The public-facing dimension of the architecture is a highly specialized mobile application that civilians install on their smartphones. This application serves a critical dual purpose: passively tracking users to assist authorities in evacuation coordination and actively allowing users to report hyper-local emergencies, such as being trapped under rubble or witnessing a medical event.[66]

Dynamic Geofencing and Background Execution Constraints

To prevent profound alert fatigue and respect user privacy, the application relies heavily on dynamic geofencing technology.[67] Utilizing the device's native GPS, cellular triangulation, and Wi-Fi data, the application establishes virtual geographic boundaries.[71] When a user crosses the threshold into a actively monitored facility—such as a large stadium, hospital, or smart building—the geofence trigger awakens the app, connecting it to the local crisis management network to receive potential alerts.[69]

Maintaining background location tracking during a crisis is technically arduous due to the

increasingly strict battery-saving protocols imposed by modern mobile operating systems. On iOS devices, the Core Location framework heavily restricts background execution; the system suspends most apps shortly after they are sent to the background to preserve battery life.[71] Furthermore, Core Location enforces hard limits, such as preventing a single application from monitoring more than 20 distinct regions simultaneously.[71] To ensure continuous emergency capability, the app must explicitly request the "Location updates" background mode capability and utilize rigorous best practices. This includes increasing the latency of location pings and utilizing larger geofence radii for areas where the user spends significant time, thereby lowering overall power consumption while still providing actionable ingress/egress data.[73]

Similarly, Android's aggressive "Doze" mode suspends network activity and CPU execution for background apps to conserve battery. To reliably wake the app and deliver critical evacuation routes or danger alerts, the backend system utilizes Firebase Cloud Messaging (FCM) payloads set exclusively to "High Priority".[75] While normal priority messages are batched and delayed indefinitely during Doze mode, high-priority FCM payloads temporarily bypass these restrictions. They grant a brief execution window for the onMessageReceived handler to process the emergency data, wake the sleeping device, and trigger a highly visible user notification, making it the only viable method for time-sensitive, life-critical alerts.[76]

Indoor Positioning Technologies

During a structural fire or building collapse, standard satellite GPS fails completely due to signal attenuation through dense concrete and steel.[79] Therefore, secondary indoor localization technologies are critical for identifying the exact coordinates of trapped users or individuals who have fainted.

The infrastructure relies on alternative indoor positioning protocols to achieve sub-room accuracy:

Technology Protocol	Expected Accuracy	Power Consumption	Implementation and Scalability Considerations
Bluetooth Low Energy (BLE / iBeacons)	2–5 meters	Very Low	Highly practical for massive, large-scale indoor deployment. Combines beacon signals with smartphone inertial sensors to track user movement

			continuously without heavy infrastructure investments.[80]
Wi-Fi FTM / Round Trip Time (RTT)	0.6–2 meters	High	Leverages existing enterprise Wi-Fi access points. Employs Conformal Prediction (CP) algorithms to provide mathematically guaranteed, uncertainty-aware prediction regions for locating trapped users.[79]
Ultra-Wideband (UWB)	10–30 cm	Medium	Expensive to deploy at scale but provides extreme, centimeter-level precision. Best reserved for highly critical or hazardous zones (e.g., chemical storage areas, complex laboratories).[80]

In an emergency, if a user activates the "Emergency Call Button" within the application, the system instantly cross-references the nearest BLE beacons or Wi-Fi RTT signals. This data is synthesized to transmit a precise 3D coordinate—including the specific floor level and room number—directly to the rescue team's operations dashboard, vastly accelerating the search and rescue timeline.[68]

Human-Centered Crisis UX/UI Design

The user experience (UX) and user interface (UI) of the mobile application are governed by strict, specialized crisis-design principles. During a disaster, a user's cognitive load spikes dramatically; they may be panicked, physically injured, or attempting to navigate through

dense, blinding smoke.[83]

Designing interfaces for these high-stress scenarios demands the ruthless elimination of all non-essential friction. Crisis UX principles dictate the following critical design metrics:

- **Prioritize Absolute Clarity Over Complexity**: Interfaces must feature minimal text, completely avoiding bureaucratic, vague, or technical jargon. Short, highly directive imperative language (e.g., "Evacuate Now" instead of "Consider evacuating if conditions worsen") is absolutely required.[84]
- **Action-Oriented Visual Hierarchy**: The most critical user actions—such as pressing the "Report Trapped" or "Medical Emergency" buttons—must utilize large, high-contrast visuals (e.g., deep red or flashing indicators) placed securely within the natural, one-handed thumb zone of the device.[83] Complex reporting tasks must be algorithmically broken down into micro-steps lasting no more than 10 to 30 seconds each, with built-in safeguards to prevent accidental, irreversible errors by shaking hands.[86]
- **Inclusion, Accessibility, and Offline Capability**: The interface must inherently support sensory-friendly notifications (for users with auditory or visual processing sensitivities), automated voice-assisted guidance, and rapid multilingual overlays to accommodate diverse urban populations.[85] Crucially, following the offline-first paradigm, the UI must gracefully degrade to local functionality, allowing users to access cached evacuation maps, progressive web app features, and vital first-aid instructions even when total network connectivity is lost.[59]

Authority Dashboards and Emergency Dispatch Integration

The vast quantities of data generated by the edge nodes, the AI triage agents, and the civilian mobile users must ultimately be synthesized into a centralized Emergency Operations Center (EOC). The EOC serves as the strategic nerve center for first responders, government authorities, and facility managers.[87]

Situational Awareness Dashboard Metrics

A successful situational awareness dashboard is defined by its ability to merge highly disparate data streams into a single, cohesive operational picture, thereby enabling rapid, data-driven decision-making.[87] Centralized, high-resolution video walls utilizing ultra-thin bezels form the core visual interface of the EOC, ergonomically positioned alongside height-adjustable workstations to reduce operator fatigue and prevent critical errors during extended, multi-day crises.[87]

The web dashboard UI dynamically overlays multiple live intelligence layers onto a core Geographic Information System (GIS) basemap:

- **Live Sensor Telemetry**: Streaming data highlighting the precise boundaries of fire spread, local seismic intensity readings, and hazardous gas concentrations.[87]
- **Predictive AI Analytics**: Visual heatmaps generated by the edge-based LSTM stampede prediction models, indicating specific corridors or intersections projected to experience crowd crush up to two minutes in the future, allowing proactive redirection.[33]
- **Live Personnel Tracking**: Real-time positional blips indicating the exact indoor coordinates of trapped civilians (via BLE/Wi-Fi RTT) alongside the current deployment vectors of police, fire, and EMS personnel.[87]

Through the implementation of these high-resolution, AI-enhanced interactive dashboards, empirical studies indicate that decision-making speed can be improved by up to 50%, transforming emergency management from a chaotic, reactive scramble into a highly orchestrated, proactive response.[87] Because this situational awareness data is highly sensitive, security is maintained through strict "secure-by-design" principles. This includes hard network isolation from standard enterprise traffic, mandatory Multi-Factor Authentication (MFA), and AI-driven continuous monitoring to detect and thwart cyber-intrusions during the vulnerable chaos of a physical disaster.[87]

NG9-1-1 Standards and PSAP Interoperability

To ensure that the localized crisis detection system operates seamlessly with municipal, state, and federal emergency services, the backend architecture must strictly adhere to the National Emergency Number Association (NENA) standards for Next Generation 9-1-1 (NG9-1-1).[90]

Legacy 9-1-1 systems, built decades ago on analog, voice-centric wireline infrastructure, are fundamentally incapable of ingesting or processing the rich multimedia data (such as live video streams, automated sensor telemetry, and precise Z-axis indoor locations) generated by modern edge AI.[92] The NENA i3 architecture standardizes the transition away from legacy circuits toward an Emergency Services IP network (ESInet), an IP-based inter-network shared by all public safety agencies.[93] A critical component of this integration is the formatting of outbound crisis data into Emergency Incident Data Objects (EIDO), which are strictly compliant with the NENA-STA-021.1b and NENA-STA-024.1.1 technical standards.[90]

When the local AI agent verifies a crisis—for instance, confirming a structural fire through visual YOLO models and thermal multimodal fusion, or detecting a localized stampede—it generates a standardized EIDO JSON payload. This payload contains the precise GIS coordinates, severity metrics, and short, critical video clips of the event.[25] This payload is transmitted directly over the secure ESInet to the Public Safety Answering Point (PSAP), completely bypassing legacy voice trunks. This advanced machine-to-machine communication ensures that 9-1-1 dispatchers receive real-time, verified, ground-truth data instantly, allowing them to route appropriate firefighting, medical, or crowd-control units without ever waiting for a panicked civilian to initiate a traditional voice call.[90]

Conclusion

The ongoing evolution of rapid crisis detection systems requires an uncompromising departure from legacy, cloud-dependent paradigms. By anchoring the architecture in high-performance edge hardware—optimizing the heavy parallel-processing strengths of the NVIDIA Jetson Orin alongside the low-power, sensory aggregation versatility of the Raspberry Pi 5—the system gains total immunity against the telecommunication decapitations that typically accompany mass disasters.

The algorithmic core of this infrastructure demonstrates that extreme accuracy is not exclusive to cloud supercomputers; it is highly achievable at the tactical edge. Utilizing sub-7-millisecond TinyML for near-field seismic anomalies, multimodal YOLO and thermal fusion for absolute fire verification, and CNN-LSTM-driven Farneback optical flow for predictive stampede mitigation, the infrastructure analyzes catastrophic threats with localized autonomy. Simultaneously, by integrating MediaPipe pose estimation and deep learning acoustic models, the architecture provides critical, life-saving detection for individual trauma, such as medical fainting or entrapment beneath building rubble.

By tying these complex hardware and software components together through robust, offline-first orchestration networks like K3s and OpenYurt, and managing the resulting data deluge through LangChain-based AI triage agents, the cognitive load on human emergency operators is drastically reduced. Finally, by extending the system's reach directly into the pockets of the public via geofenced, highly resilient mobile applications designed around crisis UX principles, and natively linking the central situational awareness dashboard to NG9-1-1 ESInet standards, the framework establishes a continuous, unbreakable chain of intelligence. From the millisecond of disaster inception to the arrival of first responders, this integrated, multi-layered approach represents the definitive engineering blueprint for future-proofing public safety in an increasingly complex and unpredictable world.

Works cited

1. Edge Computing Architecture Guide - ZPE Systems, accessed on April 11, 2026, https://zpesystems.com/edge-computing-architecture-zs/
2. Edge AI In Disaster Management - Meegle, accessed on April 11, 2026, https://www.meegle.com/en_us/topics/edge-ai-solutions/edge-ai-in-disaster-management
3. (PDF) Stampede detection and crowd analysis using CNN-LSTM ..., accessed on April 11, 2026, https://www.researchgate.net/publication/403688209_Stampede_detection_and_crowd_analysis_using_CNN-LSTM_and_farneback_optical_flow
4. Detecting Wildfire Flame and Smoke through Edge Computing using ..., accessed on April 11, 2026, https://arxiv.org/pdf/2501.08639
5. Earthquake Detection with tinyML - ResearchGate, accessed on April 11, 2026, https://www.researchgate.net/publication/370851894_Earthquake_Detection_with_tinyML

6. Raspberry Pi vs NVIDIA Jetson: the ultimate 2025 comparison - MONRASPBERRY, accessed on April 11, 2026, https://monraspberry.com/en/raspberry-pi-vs-nvidia-jetson-the-ultimate-2025-comparison/
7. Jetson Nano vs Raspberry Pi AI: The Ultimate Performance ..., accessed on April 11, 2026, https://thinkrobotics.com/blogs/learn/jetson-nano-vs-raspberry-pi-ai-the-ultimate-performance-comparison-for-edge-computing
8. Benchmarking Edge AI Platforms: Performance Analysis of NVIDIA Jetson and Raspberry Pi 5 with Coral TPU - ResearchGate, accessed on April 11, 2026, https://www.researchgate.net/publication/391165194_Benchmarking_Edge_AI_Platforms_Performance_Analysis_of_NVIDIA_Jetson_and_Raspberry_Pi_5_with_Coral_TPU
9. Raspberry Pi vs Jetson Nano: Which AI Compute Board Wins? - Zbotic, accessed on April 11, 2026, https://zbotic.in/raspberry-pi-vs-jetson-nano-which-ai-compute-board-wins/
10. (jetson nano) vs (Raspberry pi5 8g + hailo 26 TOP) - Discussion - Ultralytics, accessed on April 11, 2026, https://community.ultralytics.com/t/jetson-nano-vs-raspberry-pi5-8g-hailo-26-top/532
11. 7. Human Pose Estimation — SunFounder AI Fusion Lab Kit documentation, accessed on April 11, 2026, https://docs.sunfounder.com/projects/ai-lab-kit/en/latest/mediapipe/mp_7_pose.html
12. Uninterruptible Jetson Nano Power Supply UPS HAT Module - DFRobot, accessed on April 11, 2026, https://www.dfrobot.com/product-2342.html
13. Safe Shutdown in power outage via UPS - Raspberry Pi Forums, accessed on April 11, 2026, https://forums.raspberrypi.com/viewtopic.php?t=378626
14. Edge AI Architecture Patterns for Autonomous Systems, accessed on April 11, 2026, https://www.actian.com/blog/databases/5-edge-ai-architecture-patterns-for-disconnected-environments/
15. Raspberry Pi UPS: Best Backup Power Solutions and Setup Guide | SunFounder, accessed on April 11, 2026, https://www.sunfounder.com/blogs/news/raspberry-pi-ups-best-backup-power-solutions-and-setup-guide
16. Ups for low power devices (raspberry pi) : r/homelab - Reddit, accessed on April 11, 2026, https://www.reddit.com/r/homelab/comments/rrj612/ups_for_low_power_devices_raspberry_pi/
17. Raspberry Pi UPS? Do you need this? - YouTube, accessed on April 11, 2026, https://www.youtube.com/watch?v=oyp-m4VzxwQ
18. Survey on Reconnaissance Autonomous Robotic Systems for Disaster Management - PMC, accessed on April 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12986560/

19. Designing Energy and Power Monitoring System on Solar Power Plant Using Raspberry Pi, accessed on April 11, 2026, https://www.researchgate.net/publication/326351727_Designing_Energy_and_Power_Monitoring_System_on_Solar_Power_Plant_Using_Raspberry_Pi
20. AI-Based Smart Real-Time PV Panels Soiling Recognizing System Using Deep Neural Network Framework on NVIDIA Jetson Nano Embedded - Semantic Scholar, accessed on April 11, 2026, https://pdfs.semanticscholar.org/4be1/7009110ab9d1e27a0bc4db16e3f368ca678b.pdf
21. Edge-Computing-Enabled Smart Cities: A Comprehensive Survey - arXiv, accessed on April 11, 2026, https://arxiv.org/pdf/1909.08747
22. Earthquake detection with tinyML | U.S. Geological Survey - USGS.gov, accessed on April 11, 2026, https://www.usgs.gov/publications/earthquake-detection-tinyml
23. Early Detection of Earthquakes Using IoT and Cloud Infrastructure: A Survey - MDPI, accessed on April 11, 2026, https://www.mdpi.com/2071-1050/15/15/11713
24. A Comprehensive Review of Advanced Sensor Technologies for Fire Detection with a Focus on Gasistor-Based Sensors - MDPI, accessed on April 11, 2026, https://www.mdpi.com/2227-9040/13/7/230
25. Design and Validation of an Edge-AI Fire Safety System with SmartThings Integration for Accelerated Detection and Targeted Suppression - MDPI, accessed on April 11, 2026, https://www.mdpi.com/2076-3417/15/14/8118
26. Early Fire Detection - Multimodal AI Wildfire Prevention - MountAIn, accessed on April 11, 2026, https://www.mountainedge.ai/early-fire-detection
27. Recent Advances in Sensors for Fire Detection - PMC - NIH, accessed on April 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9100504/
28. ESCFM-YOLO: Lightweight Dual-Stream Architecture for Real-Time Small-Scale Fire Smoke Detection on Edge Devices - MDPI, accessed on April 11, 2026, https://www.mdpi.com/2076-3417/16/2/778
29. AI Fire and Smoke Detection Solution | Aetina Corporation, accessed on April 11, 2026, https://www.aetina.com/about-news-detail.php?i=1200
30. AI Fire & Smoke Detection Reimagined: Multi-Hazard Recognition from a Single Lens, accessed on April 11, 2026, https://iprogrammer.au/ai-fire-smoke-detection-reimagined/
31. Crowd Analytics & Stampede Prevention - iOmniscient, accessed on April 11, 2026, https://2025.iomni.ai/comprehensive-ai-com/crowd-analytics-stampede-prevention/
32. A Machine Learning Model for Crowd Density Classification in Hajj Video Frames - The Science and Information (SAI) Organization, accessed on April 11, 2026, https://thesai.org/Downloads/Volume15No12/Paper_31-A_Machine_Learning_Model_for_Crowd_Density_Classification.pdf
33. A Novel AI-Based Predictive Crowd Monitoring and Stampede Prevention System | Request PDF - ResearchGate, accessed on April 11, 2026, https://www.researchgate.net/publication/401007480_A_Novel_AI-Based_Predictive_Crowd_Monitoring_and_Stampede_Prevention_System

34. A Review of Abnormal Crowd Behavior Recognition Technology Based on Computer Vision, accessed on April 11, 2026, https://www.mdpi.com/2076-3417/14/21/9758
35. A Decade of You Only Look Once (YOLO) for Object Detection - arXiv, accessed on April 11, 2026, https://arxiv.org/html/2504.18586v1
36. (PDF) Robust crowd anomaly detection via hybrid ensemble learning for real-world surveillance - ResearchGate, accessed on April 11, 2026, https://www.researchgate.net/publication/397920002_Robust_crowd_anomaly_detection_via_hybrid_ensemble_learning_for_real-world_surveillance
37. Jetson 24 Pose Detection with MediaPipe - Waveshare Wiki, accessed on April 11, 2026, https://www.waveshare.com/wiki/Jetson_24_Pose_Detection_with_MediaPipe
38. Real-time Human Pose Estimation using MediaPipe | Sigmoidal, accessed on April 11, 2026, https://sigmoidal.ai/en/real-time-human-pose-estimation-using-mediapipe/
39. Reduced CPU Workload for Human Pose Detection with the Aid of a Low-Resolution Infrared Array Sensor on Embedded Systems - PMC, accessed on April 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC10708851/
40. Search and rescue: drones that detect human voices under collapsed buildings, accessed on April 11, 2026, https://hellofuture.orange.com/en/search-and-rescue-drones-that-detect-human-voices-under-collapsed-buildings/
41. AI-based Drone Assisted Human Rescue in Disaster Environments: Challenges and Opportunities - arXiv, accessed on April 11, 2026, https://arxiv.org/html/2406.15875v1
42. (PDF) Audio-Processing-Based Human Detection at Disaster Sites With Unmanned Aerial Vehicle - ResearchGate, accessed on April 11, 2026, https://www.researchgate.net/publication/341812212_Audio-Processing-Based_Human_Detection_at_Disaster_Sites_With_Unmanned_Aerial_Vehicle
43. Human Respiration and Motion Detection Based on Deep Learning and Signal Processing Techniques to Support Search and Rescue Teams - MDPI, accessed on April 11, 2026, https://www.mdpi.com/2076-3417/15/4/2097
44. Evaluation of a Sensor System for Detecting Humans Trapped under Rubble: A Pilot Study, accessed on April 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC5877370/
45. Human Detection Techniques for Search and Rescue of Trapped Victims Under Debris: A Review, accessed on April 11, 2026, https://publisher.uthm.edu.my/bookseries/index.php/eiccs/article/download/34/50/229
46. Agentic AI Aims to Cut Down Emergency Response Time When Disasters Strike | Leidos, accessed on April 11, 2026, https://www.leidos.com/insights/agentic-ai-aims-cut-down-emergency-response-time-when-disasters-strike
47. How AI tools enhance crisis management communications - Glean, accessed on April 11, 2026,

https://www.glean.com/perspectives/how-ai-tools-enhance-crisis-management-communications
48. Real-Time AI for Crisis Management: Responding Faster with Smarter Systems - Striim, accessed on April 11, 2026, https://www.striim.com/blog/real-time-ai-for-crisis-management/
49. AI Agent Framework Comparison: LangChain vs AutoGPT vs CrewAI - Fast.io, accessed on April 11, 2026, https://fast.io/resources/ai-agent-framework-comparison/
50. LangChain State of AI Agents Report: 2024 Trends, accessed on April 11, 2026, https://www.langchain.com/stateofaiagents
51. How I built a Swarm of AI Agents with LangChain | by Pranam Shetty | Medium, accessed on April 11, 2026, https://medium.com/@prxshetty/how-i-built-a-swarm-of-ai-agents-with-langchain-2e0916ce0d38
52. Autonomous AI Agents — Building Business Continuity Planning & Resilience - Medium, accessed on April 11, 2026, https://medium.com/@malcolmcfitzgerald/autonomous-ai-agents-building-business-continuity-planning-resilience-345bd9fdb949
53. LangChain AutoGPT Development - Building Autonomous AI Agents | The Uncommon Engineer, accessed on April 11, 2026, https://www.uncommonengineer.com/docs/engineer/AI/LLMLangChainProject/
54. How to Build Emergency Mental Health Detection in AI Agents - DEV Community, accessed on April 11, 2026, https://dev.to/anicca_301094325e/how-to-build-emergency-mental-health-detection-in-ai-agents-30l0
55. Enhancing AI Agents with Twilio SMS Integration, accessed on April 11, 2026, https://www.twilio.com/en-us/blog/developers/tutorials/integrations/add-sms-capabilities-ai-agent
56. Build a Multilingual AI Support Agent with SendGrid and Langchain.js | Twilio, accessed on April 11, 2026, https://www.twilio.com/en-us/blog/developers/community/build-multilingual-ai-support-agent-sendgrid-langchain
57. AI Agents: Elevating Cyber Threat Intelligence to Autonomous Response - Cyware, accessed on April 11, 2026, https://www.cyware.com/news-and-press/ai-agents-elevating-cyber-threat-intelligence-to-autonomous-response
58. Offline-First IoT: Building Resilient Devices That Work Without Internet - Metadesk Global, accessed on April 11, 2026, https://metadeskglobal.com/offline-first-iot/
59. Offline-First Architecture: Designing for Reality, Not Just the Cloud | by Jusuf Topic | Medium, accessed on April 11, 2026, https://medium.com/@jusuftopic/offline-first-architecture-designing-for-reality-not-just-the-cloud-e5fd18e50a79
60. Analysis of Offline-First App Technology in Raspberry Pi Edge Computing for Post-Disaster Hospital Situation - University of Bahrain Journals, accessed on April 11, 2026,

https://journal.uob.edu.bh/bitstreams/cd3b7bd7-2400-4d8d-9b26-562da3ca41f2/download
61. [2503.04815] Comparative Analysis of Lightweight Kubernetes Distributions for Edge Computing: Security, Resilience and Maintainability - arXiv, accessed on April 11, 2026, https://arxiv.org/abs/2503.04815
62. Cloud-Edge Orchestration for Smart Cities: A Review of Kubernetes-based Orchestration Architectures - EAI Endorsed Transactions - European Alliance for Innovation, accessed on April 11, 2026, https://publications.eai.eu/index.php/sc/article/view/1197
63. (PDF) Cloud-Edge Orchestration for Smart Cities: A Review of Kubernetes-based Orchestration Architectures - ResearchGate, accessed on April 11, 2026, https://www.researchgate.net/publication/360842690_Cloud-Edge_Orchestration_for_Smart_Cities_A_Review_of_Kubernetes-based_Orchestration_Architectures
64. Cloud-Native Workload Orchestration at the Edge: A Deployment Review and Future Directions - MDPI, accessed on April 11, 2026, https://www.mdpi.com/1424-8220/23/4/2215
65. A Resilient IoT System Architecture for Disaster Management in Collapsed Buildings, accessed on April 11, 2026, https://ieeexplore.ieee.org/document/9767132/
66. Developing real-time IoT-based public safety alert and emergency ..., accessed on April 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12334619/
67. AlertFind: Emergency and Mass Notification System, accessed on April 11, 2026, https://alertfind.com/
68. Customizable Safety Applications - Rave Mobile Safety Apps, accessed on April 11, 2026, https://www.ravemobilesafety.com/products/rave-mobile-safety-apps/
69. The Benefits of Geofence Messaging for Mass Notifications - Regroup, accessed on April 11, 2026, https://www.regroup.com/blog/the-benefits-of-geofence-messaging-for-mass-notifications/
70. Using Location-Based Safety Alerts to Protect Employees Anywhere - AlertMedia, accessed on April 11, 2026, https://www.alertmedia.com/blog/location-based-safety-alerts/
71. Geofencing iOS: Understanding the limitations - Radar, accessed on April 11, 2026, https://radar.com/blog/limitations-of-ios-geofencing
72. Using Geofencing to Actively Monitor, Collect and Share Information (with Speaker's Notes), accessed on April 11, 2026, https://tmcpfs.ops.fhwa.dot.gov/pdfs/Task-1862_Presentation_w-SpeakerNotes.pdf
73. Handling location updates in the background | Apple Developer Documentation, accessed on April 11, 2026, https://developer.apple.com/documentation/corelocation/handling-location-updates-in-the-background
74. Create and monitor geofences | Sensors and location - Android Developers, accessed on April 11, 2026, https://developer.android.com/develop/sensors-and-location/location/geofencing

75. Set the priority of a message | Firebase Cloud Messaging - Google, accessed on April 11, 2026, https://firebase.google.com/docs/cloud-messaging/customize-messages/setting-message-priority
76. Set and manage Android message priority - Firebase, accessed on April 11, 2026, https://firebase.google.com/docs/cloud-messaging/android-message-priority
77. Ensure your Firebase Cloud Messaging notifications reach your users on Android, accessed on April 11, 2026, https://firebase.blog/posts/2025/04/fcm-on-android/
78. Notifying your users with FCM - Android Developers Blog, accessed on April 11, 2026, https://android-developers.googleblog.com/2018/09/notifying-your-users-with-fcm.html
79. Robust Indoor Positioning with Hybrid WiFi RTT-RSS Signals - PMC - NIH, accessed on April 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12788237/
80. BLE vs UWB vs GPS vs WiFi: Which is the Best Indoor Positioning Technology for Personal Safety? - Seeed Studio, accessed on April 11, 2026, https://www.seeedstudio.com/blog/2025/11/13/ble-vs-uwb-vs-gps-vs-wifi-which-is-the-best-indoor-positioning-technology-for-personal-safety/
81. Indoor positioning systems compared: the practical guide for 2026 - Crowd Connected, accessed on April 11, 2026, https://www.crowdconnected.com/blog/indoor-positioning-tech-update-2026/
82. Indoor Positioning with Wi-Fi Location: A Survey of IEEE 802.11mc/az/bk Fine Timing Measurement Research - arXiv, accessed on April 11, 2026, https://arxiv.org/html/2509.03901v1
83. accessed on April 11, 2026, https://medium.com/uxcentury/ux-in-crisis-designing-for-emergency-situations-a1a970372199#:~:text=Key%20Principles%20of%20Crisis%20UX%20Design&text=Minimal%20text%2C%20large%20buttons%2C%20and,services%20or%20evacuating%20a%20building.
84. UX in Crisis: Designing for Emergency Situations | by Meri Sargsyan | UXCentury - Medium, accessed on April 11, 2026, https://medium.com/uxcentury/ux-in-crisis-designing-for-emergency-situations-a1a970372199
85. UX Design for Crisis Situations: Lessons from the Los Angeles ..., accessed on April 11, 2026, https://www.uxmatters.com/mt/archives/2025/03/ux-design-for-crisis-situations-lessons-from-the-los-angeles-wildfires.php
86. Designing For Stress And Emergency - Smashing Magazine, accessed on April 11, 2026, https://www.smashingmagazine.com/2025/11/designing-for-stress-emergency/
87. Emergency Operations Center Design Best Practices, accessed on April 11, 2026, https://www.cti.com/core-design-principles-for-eocs/
88. National Incident Management System Emergency Operations Center How-to Quick Reference Guide - FEMA, accessed on April 11, 2026, https://www.fema.gov/sites/default/files/documents/fema_eoc-quick-reference-g

uide.pdf
89. Situational Awareness Tool - Data-Smart City Solutions, accessed on April 11, 2026, https://datasmart.hks.harvard.edu/solutions/situational-awareness-tool
90. NENA Standards - National Emergency Number Association, accessed on April 11, 2026, https://www.nena.org/page/i3_stage3
91. Next Generation 911 (NG911) Standards Identification and Review - National 911 Program, accessed on April 11, 2026, https://www.911.gov/assets/National_911_Program_NG911_Standards_Identification_Analysis_2020.pdf
92. NG9-1-1 Project - National Emergency Number Association, accessed on April 11, 2026, https://www.nena.org/page/ng911_project
93. NENA i3 Standard for Next Generation 9-1-1, accessed on April 11, 2026, https://cdn.ymaws.com/www.nena.org/resource/resmgr/standards/nena-sta-010.3b-2021_i3_stan.pdf
94. Next Generation 911 (NG911) Standards Identification and Review, accessed on April 11, 2026, https://www.pa.gov/content/dam/copapwp-pagov/en/pema/documents/911/partners/documents/national-ng911-standards-identification-review.pdf